
Project Milestone: Generating Playlist Names using Tracklist Information

Team: Sofia Samaniego de la Fuente (SUID: sofiasf)

sofiasf@stanford.edu

Mentor: Richard Socher

1 Problem Description

The goal is to implement a model that generates a name for a playlist given information about its tracklist. In particular, we will attempt to construct a network that builds an internal representation of the names of the tracks included in the playlist and then uses natural language generation techniques to come up with a title that could be close to one a human might devise.

2 Data

Spotify recently released the *Million Playlist Dataset*, official website hosted at <https://recsys-challenge.spotify.com>, as part of its 2018 RecSys Challenge. It comprises a set of 1,000,000 playlist that have been created by Spotify users along with a variety of features, including playlist name, description, timestamp when the playlist was last updated, and an array of information about each track in the playlist (track name, artist, album name, duration, position in the playlist). Additionally, Spotify provides a `python` script that computes the following statistics for the dataset:

Table 1: Statistics for the *Million Playlist Dataset*

Number of playlists	1,000,000
Number of tracks	66,346,428
Number of unique tracks	2,262,292
Number of unique albums	734,684
Number of unique artists	295,860
Number of unique titles	92,944
Number of playlists with descriptions	18,760
Number of unique normalized titles	17,381
Avg playlist length	66.346428

Table 2: Top playlist names, artists, and songs (with counts)

Playlist Title	Track	Artist
Country	HUMBLE. by Kendrick Lamar	Drake
Chill	One Dance by Drake	Kanye West
Rap	Broccoli (feat. Lil Yachty) by DRAM	Kendrick Lamar
Workout	Closer by The Chainsmokers	Rihanna
Oldies	Congratulations by Post Malone	The Weeknd
Christmas	Caroline by Amin	Eminem
Rock	iSpy (feat. Lil Yachty) by KYLE	Ed Sheeran
Party	Bad and Boujee (feat. Lil Uzi Vert) by Migos	Future
Throwback	Location by Khalid	Justin Bieber
Jams	XO TOUR Llif3 by Lil Uzi Vert	J. Cole

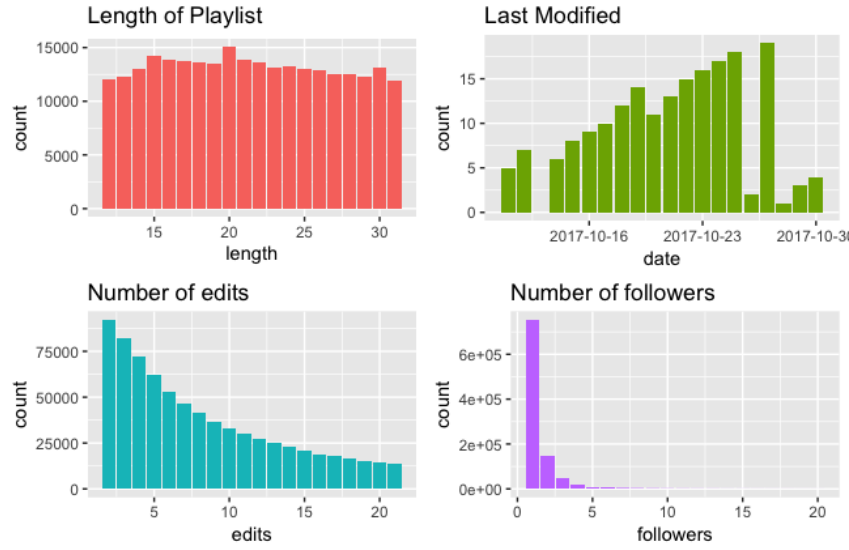


Figure 1: Histograms for selection of variables

3 Baseline

My baseline model is a sequence-to-sequence model built from scratch following Tensorflow’s Neural Machine Translation (seq2seq) Tutorial [3]. This first version is a “vanilla” NMT model, with basic unidirectional single-layer LSTM recurrent network cells with 20 units for both the encoder and the decoder. Below we show a diagram of our architecture, based on the neural machine translation architecture proposed by Luong [4].

Figure 2: Network Architecture

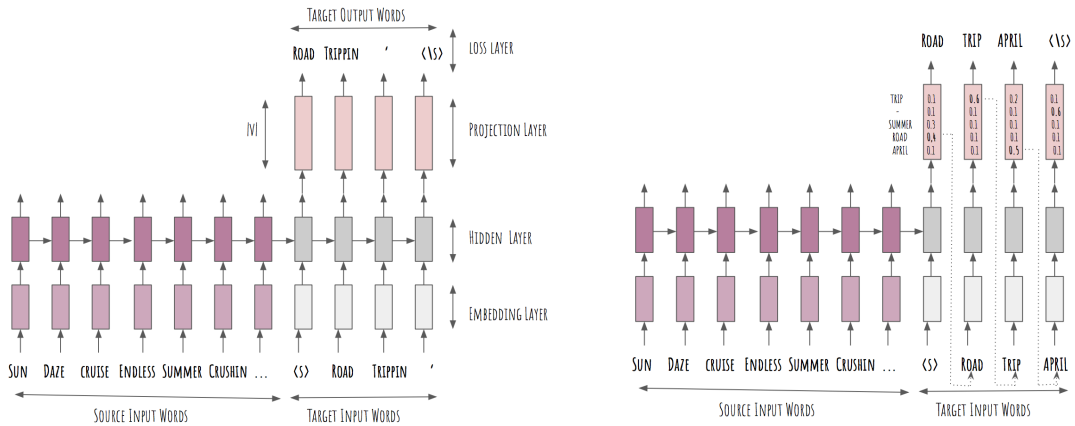


Figure 3: Train Time

Figure 4: Inference Time

Here, <s> marks the start of the decoding process while <\s> tells the decoder to stop.

- (a) **Embeddings** We initialize the embedding weights using pretrained GloVe vectors, which consist of a vocabulary of size 400,000. The embeddings are 100-dimensional. We add four additional rows to the embedding matrix, corresponding to the following special tokens:

- <unk>: Word representation for all words that appear in the track names and playlist titles but do not appear in our vocabulary.

- $\langle s \rangle$: Word representation for the start token
- $\langle \backslash s \rangle$: Word representation for the end of sentence token
- $\langle \text{mask} \rangle$: Word representation for the mask or padding token.

These four rows are initialized to zero and the embeddings from these tokens are learned from scratch during training. As a next step, we consider to restricting the GloVe vocabulary to a smaller vocabulary size V and only treating the most frequent V letters as unique, while converting every other word to the $\langle \text{unk} \rangle$ token.

- (b) **Encoder** The encoder takes the track names as inputs and does not make any predictions. Specifically, the encoder receives as input the embedded source words; that is, the word representations of the padded concatenation of all the words in each tracklist of a playlist. The input tracklists are padded to a maximum length of 250 words.
- (c) **Decoder** The decoder processes the (target) playlist names while predicting the next words. It has access to the tracklist names through the final hidden state of the encoder. Additionally, during training, the decoder receives as input the playlist names padded to a maximum length of 10 and shifted to the right by one word with an additional $\langle s \rangle$ token appended on the right. During inference, we only have access to the tracklist (source), so we cannot feed the correct (shifted) playlist names as input to the decoder. Instead, we use greedy decoding and feed the words predicted by the model in the previous timestep as input. We still use the last hidden step of the encoder to initialize the decoder, then feed a starting symbol token to the network to indicate the start of the decoding process and, in subsequent steps, feed the word with the maximum logit value out of the outputs of the decoder as input to the next timestep. The process continues until the end-of-sentence marker is produced or when we reach the maximum number of iterations.

The next steps would be to incorporate multi-layer LSTMs, add dropout, and use attention.

4 Evaluation Methodology

We will use as metric the 2-gram overlap between automatically generated playlist titles and previous titles devised by humans, as proposed in NIST's annual Document Understanding Conferences (this metric is known as ROUGE: Recall-Oriented Understudy for Gisting Evaluation). However, this could lead to bad scores for creative titles that no humans have come up with before.

5 Results

My model is currently running for 40 epochs on the full dataset on the Azure GPU. I expect it to finish by Thursday.

6 References

- [1] Cambridge Dialogue Systems Group. Rnnlg. <https://github.com/shawnwun/RNNLG>, 2017.
- [2] Facebook AI Group. Namas. <https://github.com/facebookarchive/NAMAS>, 2018.
- [3] Minh-Thang Luong, Eugene Brevdo, and Rui Zhao. Neural machine translation (seq2seq) tutorial. <https://github.com/tensorflow/nmt>, 2017.
- [4] Thang Luong. Thesis on neural machine translation. <https://github.com/lmthang/thesis>, 2017.